

Decision Trees

Sylesh Suresh

July 2017

1 Introduction

Decision trees are powerful and interpretable classifiers that mirror human decisions unlike many other classifiers in supervised machine learning and are the building blocks of random forests.

2 Definition

In essence, decision trees asks a series of true/false questions to narrow down what class a particular sample is. Here is an example of a decision tree one might use in real life to decide upon an activity on a given day:

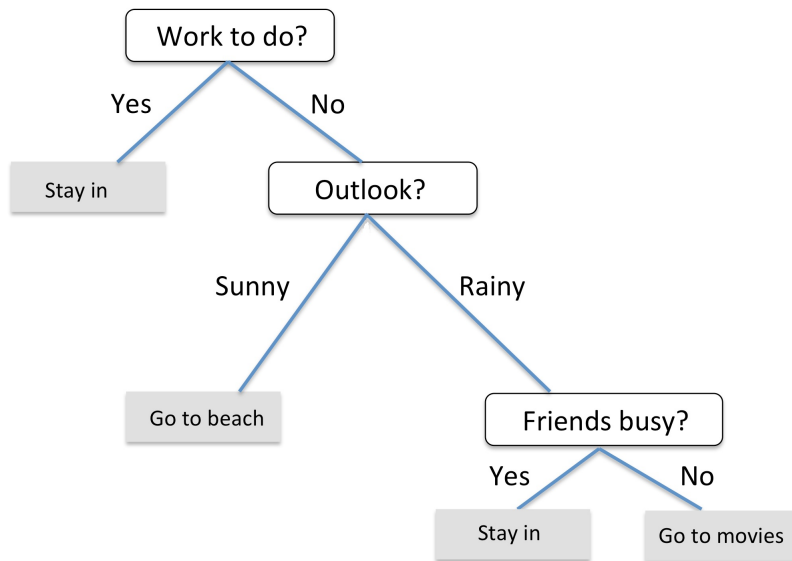


Figure 1: Real Life Decision Tree

Although this figure asks categorical variable-based questions, we can ask numerical-based questions like " $x_1 < 5$?" when the features are continuous. To build our tree, we start at the root node and ask a question that splits the data based on the feature such that the information gain is maximized. We continuously do this for each node until the decision tree can classify all the training data. Note that in practice this leads to overfitting, so the tree is usually pruned, i.e. a limit on the depth of the tree is set.

2.1 Information Gain

We split each node on the feature that yields the most information gain. The formula for information gain in a binary decision tree is as follows:

$$IG(D_p, f) = I(D_p) - \frac{N_{left}}{N_p} I(D_{left}) - \frac{N_{right}}{N_p} I(D_{right})$$

D_p is the dataset of the parent node (the node which we are splitting), f is the feature of the dataset which we are splitting on, N_p is the total number of samples in the parent node, N_{left} and N_{right} are the number of samples in the datasets of the left child node and right child node respectively, and I is the impurity measure. A node is pure if all samples in its dataset belong to the same class and is most impure when an equal number of samples belong to each class. Essentially, information gain calculates the difference between the impurity of the parent node and the impurity of the child nodes.

One commonly used measure of impurity is Gini impurity:

$$I_G(i) = 1 - \sum_{k=1}^c p(k|i)^2$$

$p(k|i)$ is the proportion of samples of class k to the total number of samples in the dataset of the i^{th} node. The impurity is maximized when the classes of the node are perfectly mixed (for this example, consider a situation in which there are 2 classes, meaning $c = 2$):

$$1 - \sum_{k=1}^c 0.5^2 = 0.5$$

An alternative impurity measure is entropy, which is defined as:

$$- \sum_{k=1}^c p(k|i) \log_2 p(k|i)$$

Note that this function has a maximum of 1.0, not 0.5. In practice, Gini impurity and entropy yield similar results, so it is more useful to test different pruning cut-offs rather than to evaluate trees with different impurity criteria.

To decide on a split for a specific node, we will search for the feature and the threshold (e.g. "petal length \leq 2.45 cm" for a flower classifier) that maximizes

the information gain. One way to choose a good threshold is to select the best threshold from the 20%, 40%, 60%, and 80% quantiles of the feature set.

An snippet from a basic implementation of a decision tree might look like this:

```
import numpy as np
def impurity(node):
    #gini impurity
    gini = 0.0
    class_values = list(set([sample[-1] for sample in node]))
    for c in class_values:
        ratio = [sample[-1] for sample in node].count(class_value) / float(len(node))
        #sample[-1] is the sample's class value

        gini += ratio*ratio

    return (1.0 - gini)

def information_gain(parent, leftchild, rightchild):
    return impurity(parent) - (len(leftchild)/len(parent))*impurity(leftchild) \
        - (len(rightchild)/len(parent))*impurity(rightchild)

def split(parent, feature, threshold):
    left = []
    right = []
    for row in parent:
        if row[feature] < threshold:
            left.append(row)
        else:
            right.append(row)
    return left, right

def best_split(parent):
    leftchild = None #placeholders
    rightchild = None
    ig = 0.0
    for feature in range(len(parent)-1): #last element is class value, not feature
        for threshold in [
            np.percentile(parent[:, feature], quartile) for quartile in [20, 40, 60, 80]]:
            if leftchild:
                p_leftchild, p_rightchild = split(parent, feature, threshold)
                p_ig = information_gain(parent, p_leftchild, p_rightchild)
                if p_ig > ig: #check if new split yields greater information gain
                    leftchild = p_leftchild
                    rightchild = p_rightchild
                    ig = p_ig
```

```
        else:
            leftchild, rightchild = split(parent, feature, threshold)
            ig = information_gain(parent, leftchild, rightchild)
    return leftchild, rightchild
```

3 Problems

1. Write a full implementation of a decision tree from scratch.
2. Write a function that returns the entropy of a node in a decision tree.